



THE PEER-TO-PEER NETWORK

JOHN NEWBERY

@jfnwbery

github.com/jnewbery

THE PEER-TO-PEER NETWORK

- ▶ Introduction
- ▶ Types of nodes
- ▶ Message format
- ▶ Control messages
- ▶ Transaction propagation
- ▶ Block propagation



THE PEER TO PEER NETWORK

WHAT IS THE PEER-TO-PEER NETWORK?

- ▶ How transactions and blocks are propagated to Bitcoin nodes
- ▶ Open, flat peer-to-peer network - no authentication, no special nodes
- ▶ Must be resistant to attacks:
 - ▶ Denial-Of-Service attacks
 - ▶ Sybil attacks

NETWORK COMMANDS

- ▶ VERSION
- ▶ VERACK
- ▶ ADDR
- ▶ GETADDR
- ▶ INV
- ▶ GETDATA
- ▶ GETBLOCKS
- ▶ GETHEADERS
- ▶ TX
- ▶ BLOCK
- ▶ HEADERS
- ▶ PING
- ▶ PONG

CONNECTING TO THE PEER-TO-PEER NETWORK

- ▶ Nodes initially connect to one or more seed nodes
- ▶ Addresses of other nodes on the network are gossiped using *ADDR* messages
- ▶ A Bitcoin Core node will connect to up to eight outbound peers
- ▶ Nodes may or may not accept inbound peers

DISCONNECTING AND BANNING (1)


- ▶ Nodes which misbehave need to be removed:
 - ▶ They waste system resources
 - ▶ They take up slots that could be used for honest peers

DISCONNECTING AND BANNING (2)

- ▶ 'bad behavior' may include:
 - ▶ Invalid transactions or blocks
 - ▶ Unconnected blocks
 - ▶ Stalling
 - ▶ *Non-standard* transactions
 - ▶ Malformed messages

DISCONNECTING AND BANNING (3)

- ▶ Depending on the misbehavior, Bitcoin Core may:
 - ▶ Ignore the problem and continue
 - ▶ Disconnect the peer immediately
 - ▶ *Ban* the peer (disconnect and don't allow connections from the same IP address for 24 hours)
 - ▶ Apply DoS points. When the DoS score reaches 100, ban the peer



TYPES OF NODES

FULL NODE

- ▶ Also called a *fully validating* node
- ▶ Receives blocks as they are mined and propagated around the network
- ▶ Verifies the validity of all blocks and all transactions included in those blocks
- ▶ Enforces the *consensus rules* of the Bitcoin network
- ▶ Maintains a collection of all the unspent outputs
- ▶ The most secure and private way to use Bitcoin

PRUNED NODE

- ▶ A type of full node
- ▶ Discards old block data to save disk space
- ▶ Retains at least 2 days of blocks and undo data to allow for re-orgs
- ▶ Propagates new blocks but cannot serve old blocks
- ▶ As secure as a non-pruned full node

'ARCHIVAL' NODE

- ▶ Unlike a pruned node, retains all old block and undo data
- ▶ Can serve old blocks to peers on the network
- ▶ Signaled using **NODE_NETWORK** in the version handshake

SIMPLE PAYMENT VERIFICATION (SPV) NODE (1)

- ▶ Only downloads:
 - ▶ the block headers
 - ▶ information about specific transactions
- ▶ Can validate proof-of-work
- ▶ Can't validate other network rules:
 - ▶ Can't detect invalid or double-spend transactions
 - ▶ Can't verify money supply

SIMPLE PAYMENT VERIFICATION (SPV) NODE (2)

- ▶ Can verify that a transaction is included in a block (by asking for Merkle proofs)
- ▶ Can't verify that a transaction hasn't appeared in the blockchain
- ▶ Can use Bloom filters to preserve (some) privacy

OTHER NODE OPTIONS

- ▶ **-blockonly** - full node which doesn't propagate transactions
- ▶ **-nolisten** - node which makes outbound connections but doesn't accept inbound connections
- ▶ **-onion** - connect to peers using Tor
- ▶ **-proxy** - connect to peers via a proxy
- ▶ **-whitelist=<IP address or subnet>** -



MESSAGE FORMAT

MESSAGE FORMAT

- ▶ Bitcoin P2P messages contain a header and a payload
- ▶ Header is 24 bytes:
 - ▶ Magic (4 bytes): indicates the network (0xf9beb4d9 for Bitcoin)
 - ▶ Command name (12 bytes): eg ADDR, INV, BLOCK, etc
 - ▶ Payload size (4 bytes): how large the payload is in bytes
 - ▶ Checksum (4 bytes): Double SHA256 of payload
- ▶ Payload: up to 32MB. Each command has its own defined format

MESSAGE FORMAT

'Magic'
bytes
(4 bytes)

Command Name
(12 bytes)

Payload
size
(4 bytes)

Checksum
(4 bytes)

Body

EXAMPLE HEADER

- ▶ **f9beb4d976657261636b0000000000000000000000005df6e0e2**
- ▶ **f9beb4d9** : network magic for Bitcoin main net
- ▶ **76657261636b0000000000000000** : VERACK with zero padding
- ▶ **00000000** : payload size is zero
- ▶ **5df6e0e2** : checksum SHA256(SHA256(""))



CONTROL MESSAGES

VERSION HANDSHAKE

- ▶ P2P connection starts with a version handshake
- ▶ Used by nodes to exchange information about themselves
- ▶ A node responds to a VERSION message with VERACK

VERSION MESSAGE (1)

- ▶ Version (4 bytes)
 - ▶ *highest version the transmitting node can connect to*
- ▶ Services (8 bytes)
 - ▶ *bitfield of services supported by the transmitting node*
- ▶ Timestamp (8 bytes)
 - ▶ *Unix timestamp of transmitting node*

VERSION MESSAGE (2)

- ▶ `addr_rcv` services (8 bytes)
 - ▶ *services supported by the receiving node*
- ▶ `addr_rcv` IP address and port (16 + 2 bytes)
 - ▶ *IPv6 address and port of receiving node*
- ▶ `addr_trans` services (8 bytes)
 - ▶ *services supported by the transmitting node (should be same as Services field)*
- ▶ `addr_trans` IP address and port (16 + 2 bytes)
 - ▶ *IPv6 address and port of transmitting node*

VERSION MESSAGE (3)

- ▶ nonce (8 bytes)
 - ▶ *random number used to detect if a node is connecting to itself*
- ▶ user_agent (compactSize + len)
 - ▶ *string indicating the software the node is using*
- ▶ start_height (4 bytes)
 - ▶ *height of the transmitting node's best blockchain*
- ▶ relay (bool - optional)
 - ▶ *indicates whether **INV** or **TX** messages should be sent to the transmitting node*

OTHER CONTROL MESSAGES

- ▶ VERACK - sent in response to a VERSION message
- ▶ ADDR - gossips connection information about other nodes
- ▶ GETADDR - requests information about other nodes
- ▶ PING/PONG - confirms connectivity
- ▶ FILTERLOAD / FILTERADD / FILTERCLEAR - sets and unsets bloom filters for SPV transaction propagation

CHECK 243840
63-1403
631
DATE March 12, 2015
\$*****2,730.00
Janku Klaparski
AUTHORIZED SIGNATURE

01536641
THE 13 SECURITY FEATURES
OF THE FEDERAL RESERVE NOTE
55543
CASHIER'S CHECK
DATE September 27, 2014
AMOUNT \$ *****2,865.00
VOID AFTER 90 DAYS
K. A. Williams
Authorized Signature

13-7882
2420
607993
00 00607993
\$2,330.00
GENERAL ELECTRIC CREDIT UNION
Paul R. Pyatt
AUTHORIZED SIGNATURE

Wells Fargo Bank, N.A.
50549
DATE 09/22/2013
\$ **2,350.00
DOLLARS
36086*

Amarillo National Bank
CAS
Notice to Purchaser
An official check may be replaced or refunded in the event it is lost, destroyed or stolen after 90 days.
PAY ONE THOUSAND SIX HUNDRED EIGHTY AND
Purchaser ROBERT WALKER
TO THE ORDER OF HALINA ZAKOWICZ
0531056 111130095

THIS DOCUMENT HAS A COLORED BACKGROUND AND MICROPRINTING IN THE BORDER
MADISON
MADISON ET CIE, INC.
8745 W 3RD ST
LOS ANGELES, CA 90048
PAY TO THE ORDER OF Halina Zakowicz
TWO-THOUSAND THREE-HUNDRED-THIRTY AND
Halina Zakowicz
MEMO PAYMENT
041057 1122000661

TO VERIFY AUTHENTICITY, SEE REVERSE SIDE
Arkansas Federal CREDIT UNION
Improving Each Member's Financial Life
P.O. Box 9
Jacksonville, AR 72078-0009
DATE: 06/20/13
CHECK NO :
PAY THE SUM OF THREE THOUSAND SEVEN HUNDRED SEVENTY-
PAY TO THE ORDER OF Halina Zakowicz
REF
575449 28207502810

TO VERIFY AUTHENTICITY, SEE REVERSE SIDE
Arkansas Federal CREDIT UNION
Improving Each Member's Financial Life
P.O. Box 9
Jacksonville, AR 72078-0009
DATE: 06/20/2013
CHECK NO :
PAY THE SUM OF THREE THOUSAND SEVEN HUNDRED SEVENTY-
PAY TO THE ORDER OF Halina Zakowicz
REF

TRANSACTION
PROPAGATION

INVENTORY ANNOUNCEMENT

- ▶ New transactions are announced in an INV message
- ▶ INV messages contain the txid
 - ▶ (Can also contain block hashes)
- ▶ If the receiving node wants the announced inventory, it responds with a GETDATA message
- ▶ The announcing node then sends a TX messages

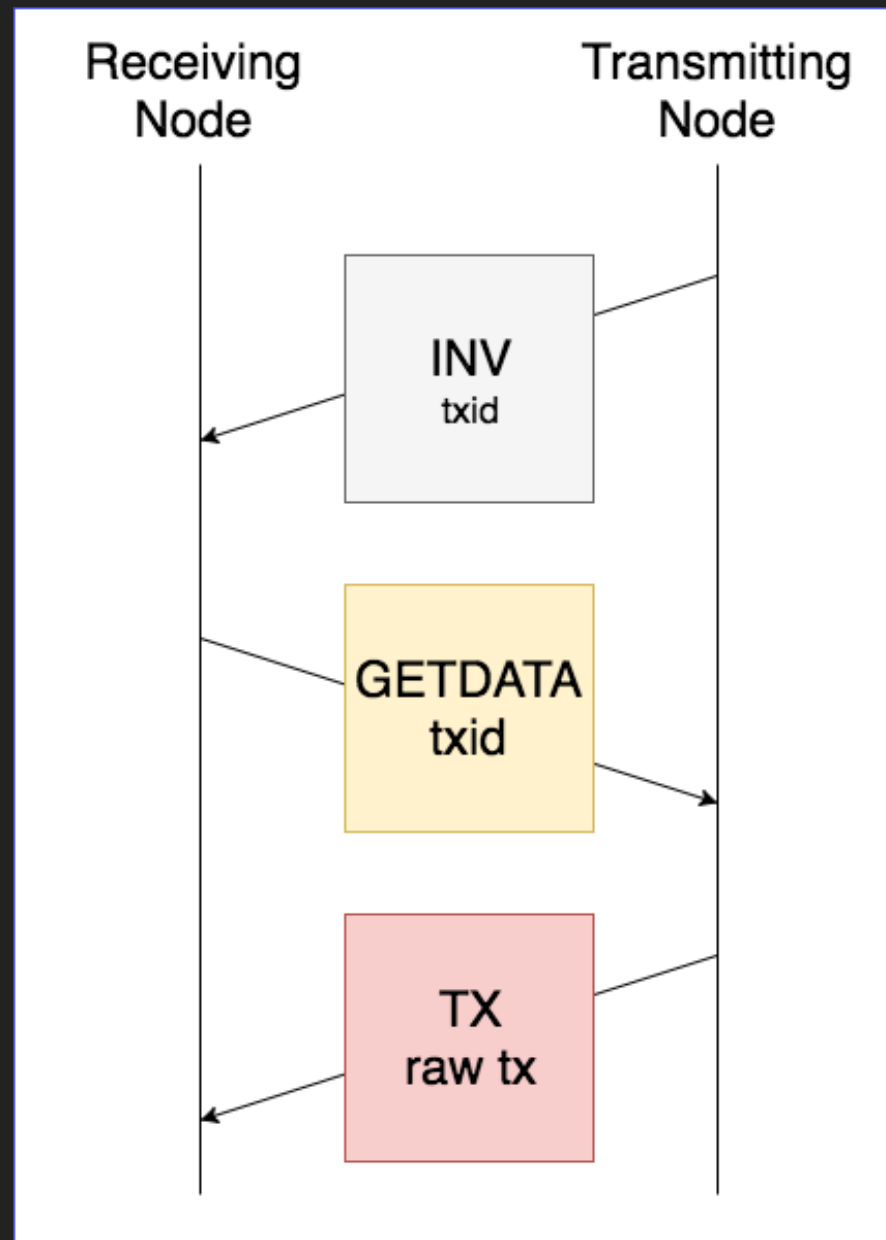
Receiving
Node


Transmitting
Node

INV
txid

GETDATA
txid

TX
raw tx





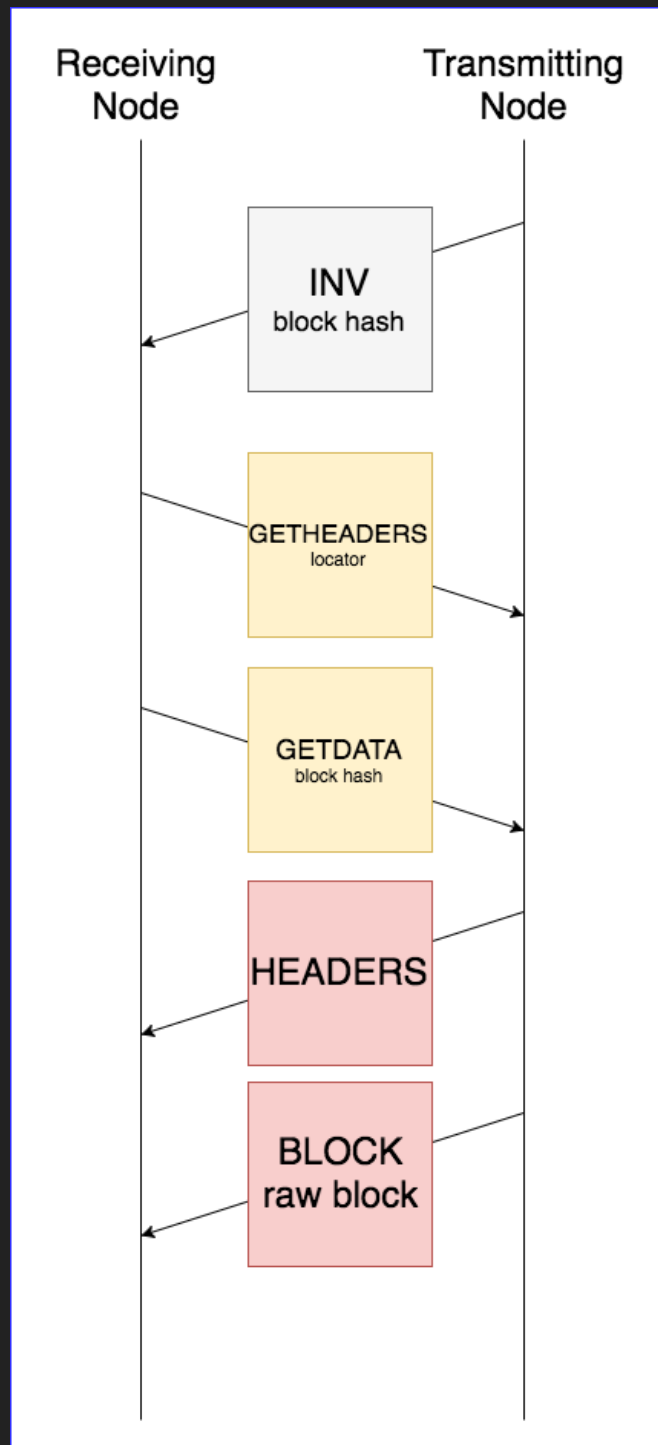
BLOCK PROPAGATION

BLOCK PROPAGATION

- ▶ Originally, blocks were propagated using INV-GETDATA-BLOCK
- ▶ v0.10.0 introduced 'headers first' syncing
- ▶ v0.12.0 introduced the SENDHEADERS message
- ▶ v0.13.0 introduced compact blocks
- ▶ v0.14.0 introduced High Bandwidth compact blocks

HEADERS-FIRST SYNCING

- ▶ Transmitting node sends INV with block hash as normal
- ▶ Receiving node responds with:
 - ▶ GETHEADERS (for block headers up to the tip)
 - ▶ GETDATA (for the tip block)
- ▶ Transmitting node sends:
 - ▶ HEADERS (connecting tip to the receiving node's best block)
 - ▶ BLOCK (containing the tip block)



SENDHEADERS

- ▶ SENDHEADERS is a new control message in protocol version 70012
- ▶ Sent immediately after VERSION handshake
- ▶ Indicates that the transmitting node would prefer to receive HEADERS messages instead of INVs
- ▶ Saves a INV-GETHEADERS round trip
- ▶ Defined in BIP 130

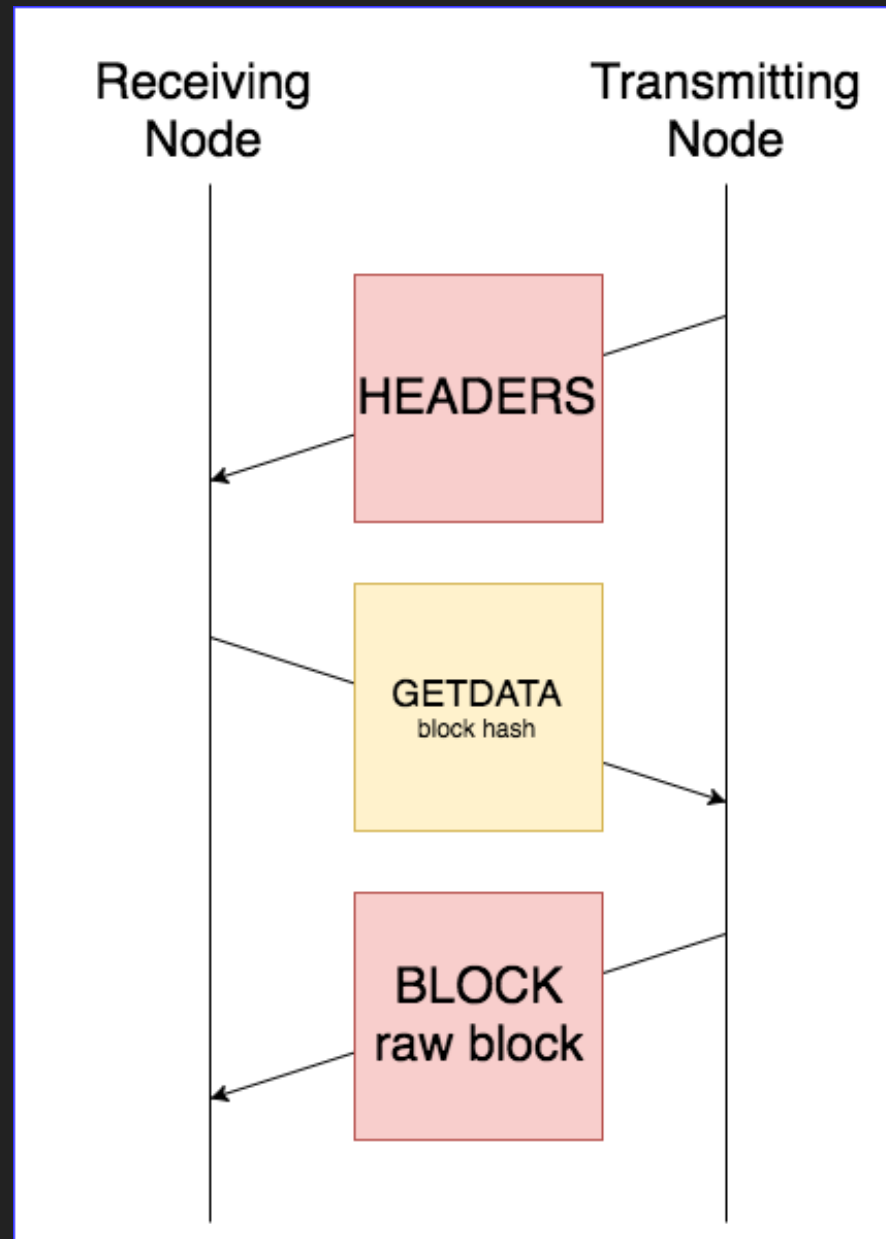
Receiving
Node

Transmitting
Node

HEADERS

GETDATA
block hash

BLOCK
raw block



COMPACT BLOCKS (1)

- ▶ Reduces time and bandwidth for propagating blocks
- ▶ Relies on fact that peer has already seen most transactions in a new block
- ▶ Enabled by node sending a SENDCMPCT message (similar to SENDHEADERS)
- ▶ Defined in BIP 152

COMPACT BLOCKS (2)

- ▶ Two modes:
 - ▶ low bandwidth - same number of messages as headers first block syncing, but saves on number of transactions sent
 - ▶ high bandwidth - sends cmpctblock message before the block has even been validated

